


8. Backup & Disaster Recovery Strategy

This solution implements layered backup and disaster recovery mechanisms across both cloud-hosted and on-premise components.

The strategy focuses on **data durability, recoverability, and operational safety**, while avoiding unnecessary over-engineering for a proof-of-concept scope.



8.1 Amazon RDS (MySQL) – Backup & Disaster Recovery (Implemented)

Amazon RDS provides native, managed backup and availability features that are leveraged to meet disaster recovery requirements with minimal operational overhead.

Automated Backups and Point-in-Time Recovery (PITR)

```
backup_retention_period = 7
```

Behavior

- AWS automatically takes daily snapshots of the database
- Transaction logs are continuously captured
- Enables point-in-time recovery within the last seven days

Design Value

- Database can be restored to any second within the retention window
 - Protects against accidental deletes, faulty deployments, and data corruption
 - No manual snapshot scheduling required
-

Multi-AZ Deployment (High Availability)

```
multi_az = true
```

Behavior

- A synchronous standby replica is created in a separate Availability Zone
- AWS automatically manages replication and failover

Design Value

- Automatic failover during AZ outages
- Minimal application downtime
- Provides availability and resilience (not read scaling)

Note: Multi-AZ is an availability mechanism, not a scaling strategy.

Final Snapshot on Deletion (Data Safety Net)

```
skip_final_snapshot      = false
final_snapshot_identifier = "bookdb-final-snapshot"
```

Behavior

- A final snapshot is taken before the RDS instance is deleted
- Snapshot is retained independently of the database lifecycle

Design Value

- Prevents irreversible data loss
 - Enables full database restoration after accidental deletion
 - Acts as a last-resort recovery mechanism
-

Deletion Protection (Human Error Mitigation)

```
deletion_protection = true
```

Behavior

- Prevents accidental deletion via Terraform or AWS Console
- Deletion requires explicit disabling of protection

Design Value

- Safeguards production data
 - Enforces intentional, reviewed destructive actions
-

Storage Encryption (Data at Rest)

```
storage_encrypted = true
```

Behavior

- Encrypts database storage
- Encrypts automated backups and snapshots

Design Value

- Aligns with security best practices and compliance expectations
 - Protects data even if storage media is accessed improperly
-

Maintenance Window and Controlled Patching

```
maintenance_window      = "sun:02:00-sun:03:00"  
auto_minor_version_upgrade = true
```

Behavior

- AWS applies minor version upgrades and patches only during the defined window

Design Value

- Predictable maintenance behavior
 - Avoids unplanned downtime during business hours
 - Reduces operational surprises
-

Snapshot Tagging (Operational Clarity)

```
copy_tags_to_snapshot = true
```

Behavior

- Automatically copies resource tags to snapshots

Design Value

- Simplifies snapshot identification
- Improves cost allocation and auditability
- Supports cleaner operations in multi-environment setups

Restore Capabilities (Disaster Scenarios)

With this configuration, the database can be:

- Restored to a new RDS instance
- Restored into a different Availability Zone
- Copied and restored into another region or VPC (manual cross-region copy)

This satisfies disaster recovery requirements without introducing unnecessary complexity.

8.2 On-Premise Backup Strategy (Conceptual, Realistic)

Unlike Amazon RDS, on-premise systems do not provide managed backups. Backups are therefore explicitly automated and stored off-site in Amazon S3.

Backup Scope

The following components are backed up:

- MySQL database (`mysqldump`)
- Application artifacts
- Configuration files

Backup Flow (Conceptual)

On-Prem Server



`mysqldump / tar`



`Upload to Amazon S3`



`S3 Lifecycle Policies`

This approach ensures off-site, durable backups with no dependency on local storage.

Example Backup Script (Conceptual)

```
#!/bin/bash
mysqldump -u root -p appdb > appdb.sql
aws s3 cp appdb.sql s3://onprem-backups/appdb-$(date +%F).sql
```

Design Value

- Enables regular, automated off-site backups
 - Protects against on-prem hardware failure
 - Allows straightforward restoration into Amazon RDS if required
-

S3 Lifecycle Policies (Retention and Cost Control)

Configured conceptually to:

- Retain daily backups for 30 days
- Transition older backups to lower-cost storage (e.g., Glacier)
- Automatically expire outdated backups

Why Amazon S3

- Extremely high durability (11 nines)
 - Cost-effective compared to block storage
 - Native lifecycle management simplifies retention enforcement
-

8.3 Design Intent Summary

- Cloud-hosted databases rely on managed AWS backup and availability features

- On-premise systems use scripted backups with off-site storage
- Data protection is layered, intentional, and operationally simple
- The strategy balances resilience with PoC-appropriate complexity